



ARUNAI ENGINEERING COLLEGE

(An Autonomous Institution)
Velu Nagar, Thiruvannamalai-606603
www.arunai.org



DEPARTMENT OF ARTIFICIAL INTELLIGENCE & DATA SCIENCE

BACHELOR OF TECHNOLOGY

2024-2025

FOURTH SEMESTER

**AD3461-Machine Learning
Laboratory**

ARUNAI ENGINEERING COLLEGE

TIRUVANNAMALAI – 606 603



DEPARTMENT OF ARTIFICIAL INTELLIGENCE & DATA SCIENCE CERTIFICATE

Certified that this is a bonafide record of work done by

Name :

University Reg.No :

Semester :

Branch :

Year :

Staff-in-Charge

Head of the Department

Submitted for the _____

Practical Examination held on _____

Internal Examiner

External Examiner

S.NO	Date	Name of Experiments	Page No	Signature
1		Candidate-Elimination algorithm		
2		Decision tree based ID3 algorithm		
3		Artificial Neural Network by the Backpropagation algorithm		
4		Naive Bayesian classifier to compute the accuracy with a few test data sets		
5		Naive Bayesian Classifier to compute the accuracy, precision, and recall.		
6		Bayesian network to diagnose CORONA infection		
7		EM algorithm to cluster a set of data using the k-Means algorithm		
8		K-Nearest Neighbour algorithm to classify the iris data set.		
9		Locally Weighted Regression algorithm to fit data points		
10		K-Fold validation		
11		Paired T – test		

PROGRAM:

```
import numpy as np
import pandas as pd
data = pd.read_csv(path+'/enjoysport.csv')
concepts = np.array(data.iloc[:,0:-1])
print("\nInstances are:\n",concepts)
target = np.array(data.iloc[:,-1])
print("\nTarget Values are: ",target)

def learn(concepts, target):
    specific_h = concepts[0].copy()
    print("\nInitialization of specific_h and general_h")
    print("\nSpecific Boundary: ", specific_h)
    general_h = [["?" for i in range(len(specific_h))] for i in range(len(specific_h))]
    print("\nGeneric Boundary: ",general_h)

    for i, h in enumerate(concepts):
        print("\nInstance", i+1 , "is ", h)
        if target[i] == "yes":
            print("Instance is Positive ")
            for x in range(len(specific_h)):
                if h[x] != specific_h[x]:
                    specific_h[x] = '?'
                    general_h[x][x] = '?'

        if target[i] == "no":
            print("Instance is Negative ")
            for x in range(len(specific_h)):
                if h[x] != specific_h[x]:
                    general_h[x][x] = specific_h[x]
            else:
                general_h[x][x] = '?'
```

```
print("Specific Bunday after ", i+1, "Instance is ", specific_h)
print("Generic Boundary after ", i+1, "Instance is ", general_h)
print("\n")
```

```
indices = [i for i, val in enumerate(general_h) if val == ['?', '?', '?', '?', '?', '?']] for i in
indices:
```

```
general_h.remove(['?', '?', '?', '?', '?', '?']) return
```

```
specific_h, general_h
```

```
s_final, g_final = learn(concepts, target)
```

```
print("Final Specific_h: ", s_final, sep="\n")
```

```
print("Final General_h: ", g_final, sep="\n")
```

DATASET:

	Outlook	Temperature	Humidity	Wind	Answer
1	sunny	hot	high	weak	No
2	sunny	hot	high	strong	No
3	overcast	hot	high	weak	Yes
4	rain	mild	high	weak	Yes
5	rain	cool	normal	weak	Yes
6	rain	cool	normal	strong	No
7	overcast	cool	normal	strong	Yes
8	sunny	mild	high	weak	No
9	sunny	cool	normal	weak	Yes
10	rain	mild	normal	weak	Yes
11	sunny	mild	normal	strong	Yes
12	overcast	mild	high	strong	Yes
13	overcast	hot	normal	weak	Yes
14	rain	mild	high	strong	No
15	sunny	hot	high	strong	No

Output

[[‘sunny’ ‘warm’ ‘normal’ ‘strong’ ‘warm’ ‘same’]]

['sunny' 'warm' 'high' 'strong' 'warm' 'same']

['rainy' 'cold' 'high' 'strong' 'warm' 'change']

[‘sunny’ ‘warm’ ‘high’ ‘strong’ ‘cool’ ‘change’]]

Target Values are: ['yes' 'yes' 'no' 'yes']

Initialization of specific_h and general_h

Specific Boundary: ['sunny' 'warm' 'normal' 'strong' 'warm' 'same']

eneric Boundary: [[‘?', ‘?', ‘?', ‘?', ‘?', ‘?’], [‘?', ‘?', ‘?', ‘?', ‘?', ‘?’], [‘?', ‘?', ‘?', ‘?', ‘?', ‘?’], [‘?', ‘?', ‘?', ‘?', ‘?', ‘?’], [‘?', ‘?', ‘?', ‘?', ‘?', ‘?’], [‘?', ‘?', ‘?', ‘?', ‘?', ‘?’, ‘?’]]

Instance 1 is ['sunny' 'warm' 'normal' 'strong' 'warm' 'same'] Instance is

Positive

Specific Boundary after 1 Instance is ['sunny' 'warm' 'normal' 'strong' 'warm' 'same']

Generic Boundary after 1 Instance is [[‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’]]

Instance 2 is [‘sunny’ ‘warm’ ‘high’ ‘strong’ ‘warm’ ‘same’] Instance is Positive
Specific Boundary after 2 Instance is [‘sunny’ ‘warm’ ‘?’ ‘strong’ ‘warm’ ‘same’]
Generic Boundary after 2 Instance is [[‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’]]

Instance 3 is [‘rainy’ ‘cold’ ‘high’ ‘strong’ ‘warm’ ‘change’] Instance is Negative
Specific Boundary after 3 Instance is [‘sunny’ ‘warm’ ‘?’ ‘strong’ ‘warm’ ‘same’]
Generic Boundary after 3 Instance is [[‘sunny’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘warm’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘same’]]

Instance 4 is [‘sunny’ ‘warm’ ‘high’ ‘strong’ ‘cool’ ‘change’] Instance is Positive
Specific Boundary after 4 Instance is [‘sunny’ ‘warm’ ‘?’ ‘strong’ ‘?’ ‘?’]
Generic Boundary after 4 Instance is [[‘sunny’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘warm’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’]]

Final Specific_h: [‘sunny’ ‘warm’ ‘?’ ‘strong’ ‘?’ ‘?’]

Final General_h: [[‘sunny’ , ‘?’ , ‘?’ , ‘?’ , ‘?’ , ‘?’], [‘?’ , ‘warm’ , ‘?’ , ‘?’ , ‘?’ , ‘?’]]

PROGRAM:

```
import pandas as pd
import math
import numpy as np

data = pd.read_csv("Datasets/enjoysport.csv")
features = [feat for feat in data.columns if feat != "Answer"] # Filter out 'Answer' column

class Node:
    def __init__(self):
        self.children = []
        self.value = ""
        self.isLeaf = False
        self.pred = ""

def entropy(examples):
    pos = 0.0
    neg = 0.0
    for _, row in examples.iterrows():
        if row["Answer"] == "yes":
            pos += 1
        else:
            neg += 1
    if pos == 0.0 or neg == 0.0:
        return 0.0
    else:
        p = pos / (pos + neg)
        n = neg / (pos + neg)
        return -(p * math.log(p, 2) + n * math.log(n, 2))

def info_gain(examples, attr):
    uniq = np.unique(examples[attr])
    gain = entropy(examples)
    for u in uniq:
        subdata = examples[examples[attr] == u]
        sub_e = entropy(subdata)
        gain -= (float(len(subdata)) / float(len(examples))) * sub_e
    return gain

def ID3(examples, attrs):
```



```

root = Node()
max_gain = 0
max_feat = ""
for feature in attrs:
    gain = info_gain(examples, feature)
    if gain > max_gain:
        max_gain = gain
        max_feat = feature

if not max_feat:
    root.isLeaf = True
    root.pred = examples["Answer"].mode()[0]
    return root

root.value = max_feat
uniq = np.unique(examples[max_feat])
for u in uniq:
    subdata = examples[examples[max_feat] == u]
    if entropy(subdata) == 0.0:
        newNode = Node()
        newNode.isLeaf = True
        newNode.value = u
        newNode.pred = np.unique(subdata["Answer"])
        root.children.append(newNode)
    else:
        dummyNode = Node()
        dummyNode.value = u
        new_attrs = attrs.copy()
        new_attrs.remove(max_feat)
        child = ID3(subdata, new_attrs)
        dummyNode.children.append(child)
        root.children.append(dummyNode)
return root

```

```

def printTree(root: Node, depth=0):
    for i in range(depth):
        print("\t", end="")
    print(root.value, end="")
    if root.isLeaf:
        print(" -> ", root.pred)

```

```

    print()
    for child in root.children:
        printTree(child, depth + 1)

def classify(root: Node, new):
    for child in root.children:
        if child.value == new[root.value]:
            if child.isLeaf:
                print("Predicted Label for new example", new, " is:", child.pred)
                return
            else:
                classify(child.children[0], new)

root = ID3(data, features)
print("Decision Tree is:")
printTree(root)
print(" -----")
new = {"Outlook": "sunny", "Temperature": "hot", "Humidity": "normal", "Wind": "strong"}
classify(root, new

```

DATASET:

	Outlook	Temperature	Humidity	Wind	Answer
1	sunny	hot	high	weak	No
2	sunny	hot	high	strong	No
3	overcast	hot	high	weak	Yes
4	rain	mild	high	weak	Yes
5	rain	cool	normal	weak	Yes
6	rain	cool	normal	strong	No
7	overcast	cool	normal	strong	Yes
8	sunny	mild	high	weak	No
9	sunny	cool	normal	weak	Yes
10	rain	mild	normal	weak	Yes
11	sunny	mild	normal	strong	Yes
12	overcast	mild	high	strong	Yes
13	overcast	hot	normal	weak	Yes
14	rain	mild	high	strong	No
15	sunny	hot	high	strong	No

OUTPUT

Decision Tree is:

Outlook

overcast -> ['yes']

rain

Wind

strong -> ['no']

weak -> ['yes']

sunny

Humidity

high -> ['no']

normal -> ['yes']

Predicted Label for new example {'Outlook': 'sunny', 'Temperature': 'hot', 'Humidity': 'normal', 'Wind': 'strong'} is: ['yes']

PROGRAM:

```
import numpy as np

X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float)

y = np.array([[92], [86], [89]], dtype=float)

X = X/np.amax(X,axis=0) #maximum of X array longitudinally

y = y/100
```

#Sigmoid Function

```
def sigmoid (x):

    return 1/(1 + np.exp(-x))
```

#Derivative of Sigmoid Function

```
def derivatives_sigmoid(x):

    return x * (1 - x)
```

#Variable initialization

```
epoch=5 #Setting training iterations

lr=0.1 #Setting learning rate

inputlayer_neurons = 2 #number of features in data set

hiddenlayer_neurons = 3 #number of hidden layers neurons

output_neurons = 1 #number of neurons at output layer
```

#weight and bias initialization

```
wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))

bh=np.random.uniform(size=(1,hiddenlayer_neurons))

wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
```

```

bout=np.random.uniform(size=(1,output_neurons))

#draws a random range of numbers uniformly of dim x*y

for i in range(epoch):

    #Forward Propogation

    hinp1=np.dot(X,wh)

    hinp=hinp1 + bh

    hlayer_act = sigmoid(hinp)

    outinp1=np.dot(hlayer_act,wout)

    outinp= outinp1+bout

    output = sigmoid(outinp)

#Backpropagation

    EO = y-output

    outgrad = derivatives_sigmoid(output)

    d_output = EO * outgrad

    EH = d_output.dot(wout.T)

    hiddengrad = derivatives_sigmoid(hlayer_act)

#how much hidden layer wts contributed to error

    d_hiddenlayer = EH * hiddengrad

    wout += hlayer_act.T.dot(d_output) *lr

# dot product of next layer error and current layerop

    wh += X.T.dot(d_hiddenlayer) *lr

    print ("-----Epoch-", i+1, "Starts ----- ")

    print("Input: \n" + str(X))

```

```
print("Actual Output: \n" + str(y))  
print("Predicted Output: \n",output)  
print ("-----Epoch-", i+1, "Ends ----- \n")  
print("Input: \n" + str(X))  
print("Actual Output: \n" + str(y))  
print("Predicted Output: \n",output)
```

Training Examples:

Example	Sleep	Study	Expected % in Exams
1	2	9	92
2	1	5	86
3	3	6	89

Normalize the input

Example	Sleep	Study	Expected % in Exams
1	$2/3 = 0.66666667$	$9/9 = 1$	0.92
2	$1/3 = 0.33333333$	$5/9 = 0.55555556$	0.86
3	$3/3 = 1$	$6/9 = 0.66666667$	0.89

Output

————Epoch- 1 Starts————

Input:

[[0.66666667 1.]

[0.33333333 0.55555556]

[1. 0.66666667]]

Actual Output:

[[0.92]

[0.86]

[0.89]]

Predicted Output:

[[0.81951208]

[0.8007242]

[0.82485744]]

————Epoch- 1 Ends————

PROGRAM:

```
import pandas as pd
from sklearn import tree
from sklearn.preprocessing import LabelEncoder
from sklearn.naive_bayes import GaussianNB

# load data from CSV
data = pd.read_csv('tennisdata.csv')
print("The first 5 values of data is :\n",data.head())

# obtain Train data and Train output
X = data.iloc[:, :-1]
print("\nThe First 5 values of train data is\n",X.head())

y = data.iloc[:, -1]
print("\nThe first 5 values of Train output is\n",y.head())

# Convert then in numbers
le_outlook = LabelEncoder()
X.Outlook = le_outlook.fit_transform(X.Outlook)

le_Temperature = LabelEncoder()
X.Temperature = le_Temperature.fit_transform(X.Temperature)

le_Humidity = LabelEncoder()
X.Humidity = le_Humidity.fit_transform(X.Humidity)

le_Windy = LabelEncoder()
X.Windy = le_Windy.fit_transform(X.Windy)

print("\nNow the Train data is :\n",X.head())

le_PlayTennis = LabelEncoder()
y = le_PlayTennis.fit_transform(y)
```

```
print("\nNow the Train output is\n",y)
```

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.20)
```

```
classifier = GaussianNB()  
classifier.fit(X_train,y_train)
```

```
from sklearn.metrics import accuracy_score  
print("Accuracy is:",accuracy_score(classifier.predict(X_test),y_test))
```

Tennisdata.Csv

Outlook	Temperature	Humidity	Windy	Play Tennis
Sunny	Hot	High	FALSE	No
Sunny	Hot	High	TRUE	No
Overcast	Hot	High	FALSE	Yes
Rainy	Mild	High	FALSE	Yes
Rainy	Cool	Normal	FALSE	Yes
Rainy	Cool	Normal	TRUE	No
Overcast	Cool	Normal	TRUE	Yes
Sunny	Mild	High	FALSE	No
Sunny	Cool	Normal	FALSE	Yes
Rainy	Mild	Normal	FALSE	Yes
Sunny	Mild	Normal	TRUE	Yes
Overcast	Mild	High	TRUE	Yes
Overcast	Hot	Normal	FALSE	Yes

OUTPUT:

The first 5 values of data is :

	Outlook	Temperature	Humidity	Windy	PlayTennis
0	Sunny	Hot	High	False	No
1	Sunny	Hot	High	True	No
2	Overcast	Hot	High	False	Yes
3	Rainy	Mild	High	False	Yes
4	Rainy	Cool	Normal	False	Yes

The First 5 values of train data is

	Outlook	Temperature	Humidity	Windy
0	Sunny	Hot	High	False
1	Sunny	Hot	High	True
2	Overcast	Hot	High	False
3	Rainy	Mild	High	False
4	Rainy	Cool	Normal	False

The first 5 values of Train output is

0	No
1	No
2	Yes
3	Yes
4	Yes

Name: PlayTennis, dtype: object

Now the Train data is :

	Outlook	Temperature	Humidity	Windy
0	2	1	0	0
1	2	1	0	1
2	0	1	0	0
3	1	2	0	0
4	1	0	1	0

Now the Train output is

[0 0 1 1 1 0 1 0 1 1 1 1 1 0]

Accuracy is: 0.6666666666666666



PROGRAM:

```
import numpy as np
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, precision_score, recall_score

# Sample documents and their corresponding labels
documents = [
    "This is a sample document about the naive Bayes classifier algorithm.",
    "Naive Bayes classifier is easy to implement and works well for text classification tasks.",
    "Text classification using the naive Bayes algorithm is popular in natural language processing.",
    "The output of the naive Bayes program depends on the input features and training data."
]
labels = [1, 1, 1, 0] # 1 for documents about naive Bayes, 0 for others

# Convert the documents into a bag-of-words representation
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(documents)

# Train a naive Bayes classifier
classifier = MultinomialNB()
classifier.fit(X, labels)

# Test data
test_documents = [
    "This document is not related to naive Bayes.",
    "Naive Bayes algorithm is widely used for text classification."
]
true_labels = [0, 1]

# Convert test documents into bag-of-words representation
X_test = vectorizer.transform(test_documents)

# Predict labels for test documents
predicted_labels = classifier.predict(X_test)

# Calculate accuracy, precision, and recall
accuracy = accuracy_score(true_labels, predicted_labels)
precision = precision_score(true_labels, predicted_labels)
```

```
recall = recall_score(true_labels, predicted_labels)
```

```
print("Accuracy:", accuracy)
```

```
print("Precision:", precision)
```

```
print("Recall:", recall)
```

OUTPUT:

Accuracy: 0.5

Precision: 0.5

Recall: 1.0

PROGRAM:

```
# Define the conditional probabilities
# P(COVID | Fever, Cough, BreathingDifficulty)
P_COVID_given_symptoms = {
    (1, 1, 1): 0.15, # P(COVID = Yes | Fever = Yes, Cough = Yes, BreathingDifficulty = Yes)
    (1, 1, 0): 0.1, # P(COVID = Yes | Fever = Yes, Cough = Yes, BreathingDifficulty = No)
    (1, 0, 1): 0.4, # P(COVID = Yes | Fever = Yes, Cough = No, BreathingDifficulty = Yes)
    (1, 0, 0): 0.1, # P(COVID = Yes | Fever = Yes, Cough = No, BreathingDifficulty = No)
    (0, 1, 1): 0.7, # P(COVID = Yes | Fever = No, Cough = Yes, BreathingDifficulty = Yes)
    (0, 1, 0): 0.3, # P(COVID = Yes | Fever = No, Cough = Yes, BreathingDifficulty = No)
    (0, 0, 1): 0.2, # P(COVID = Yes | Fever = No, Cough = No, BreathingDifficulty = Yes)
    (0, 0, 0): 0.01 # P(COVID = Yes | Fever = No, Cough = No, BreathingDifficulty = No)
}

# Given evidence
fever = 1 # Fever = Yes
cough = 1 # Cough = Yes
breathing_difficulty = 1 # BreathingDifficulty = Yes

# Calculate P(COVID = Yes) and P(COVID = No)
P_COVID_yes = P_COVID_given_symptoms[(fever, cough, breathing_difficulty)]
P_COVID_no = 1 - P_COVID_yes

# Print the result in the desired format
print("COVID P(COVID)")
print(f"0 {P_COVID_no:.2f}")
print(f"1 {P_COVID_yes:.2f}")
```

OUTPUT:

COVID	P(COVID)
0	0.85
1	0.15

PROGRAM:

```
import numpy as np
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt

data = pd.read_csv('dataset1.csv')
X = data.values

num_clusters = 2

kmeans = KMeans(n_clusters=num_clusters, random_state=42)
kmeans_labels = kmeans.fit_predict(X)
kmeans_silhouette_score = silhouette_score(X, kmeans_labels)

em = GaussianMixture(n_components=num_clusters, random_state=42)
em_labels = em.fit_predict(X)
em_silhouette_score = silhouette_score(X, em_labels)

print("Silhouette Score (k-Means):", kmeans_silhouette_score)
print("Silhouette Score (EM):", em_silhouette_score)

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.scatter(X[:, 0], X[:, 1], c=kmeans_labels, cmap='viridis')
plt.scatter(kmeans.cluster_centers[:, 0], kmeans.cluster_centers[:, 1], marker='x', color='red',
label='Centroids')
plt.title('k-Means Clustering')
plt.legend()

plt.subplot(1, 2, 2)
plt.scatter(X[:, 0], X[:, 1], c=em_labels, cmap='viridis')
plt.scatter(em.means[:, 0], em.means[:, 1], marker='x', color='red', label='Centroids')
plt.title('EM Clustering')
plt.legend()
```

```
plt.show()
```

DATASET

Feature1,Feature2

2.5,3.5

1.5,2.5

3.5,4.5

3.0,4.0

2.0,3.0

7.5,6.5

8.5,7.5

9.0,8.0

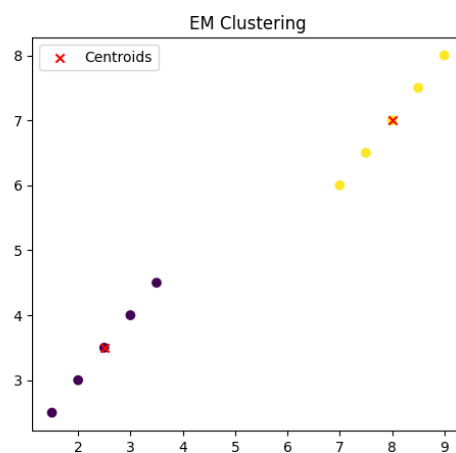
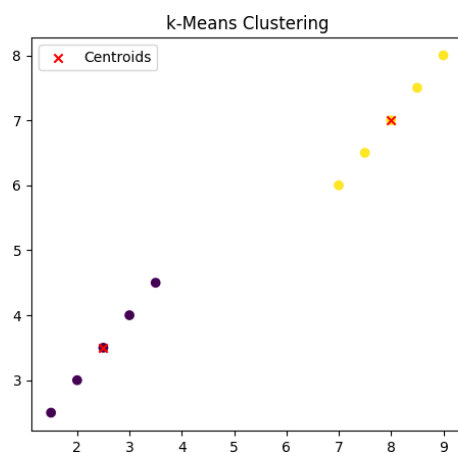
8.0,7.0

7.0,6.0

OUTPUT:

Silhouette Score (k-Means): 0.7774804461410134

Silhouette Score (EM): 0.7774804461410134



PROGRAM:

```
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

iris = load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

k = 3

knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

for i in range(len(X_test)):
    if y_pred[i] == y_test[i]:
        print(f'Correct prediction: Actual - {iris.target_names[y_test[i]]}, Predicted - {iris.target_names[y_pred[i]]}')
    else:
        print(f'Wrong prediction: Actual - {iris.target_names[y_test[i]]}, Predicted - {iris.target_names[y_pred[i]]}')

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

[illegible]

Correct prediction: Actual - setosa, Predicted - setosa

Correct prediction: Actual - virginica, Predicted - virginica

Correct prediction: Actual - versicolor, Predicted - versicolor

Correct prediction: Actual - versicolor, Predicted - versicolor

Correct prediction: Actual - setosa, Predicted - setosa

Correct prediction: Actual - setosa, Predicted - setosa

Accuracy: 1.0

PROGRAM:

```
import numpy as np
import matplotlib.pyplot as plt

def lwr(query_point, X, y, tau):
    """
    Locally Weighted Regression
    Args:
    - query_point: point at which prediction is to be made
    - X: input features
    - y: target values
    - tau: bandwidth parameter
    Returns:
    - prediction at query_point
    """
    m = X.shape[0]
    X = np.column_stack((np.ones(m), X)) # Add bias term
    query_point = np.array([1, query_point]) # Add bias term to query point
    weights = np.exp(-((X[:, 1] - query_point[1])**2) / (2 * tau * tau))
    W = np.diag(weights)
    theta = np.linalg.inv(X.T @ W @ X) @ (X.T @ (W @ y))
    prediction = query_point @ theta
    return prediction


np.random.seed(0)
X = np.linspace(0, 10, 100)
y = np.sin(X) + np.random.normal(0, 0.1, 100)

tau = 1.0

predictions = [lwr(x, X, y, tau) for x in X]

plt.figure(figsize=(10, 6))
plt.scatter(X, y, color='blue', label='Original Data')
plt.plot(X, predictions, color='red', label='Fitted Curve')
plt.xlabel('X')
plt.ylabel('y')
plt.title('Locally Weighted Regression')
plt.legend()
plt.show()
```


OUTPUT:



PROGRAM:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import KFold

data = load_iris()
X = data.data
y = data.target

model = KNeighborsClassifier(n_neighbors=3)

kf = KFold(n_splits=5, shuffle=True, random_state=42)

scores = cross_val_score(model, X, y, cv=kf, scoring='accuracy')

print(f"Accuracy for each fold: {scores}")
print(f"Mean accuracy: {scores.mean():.2f}")
print(f"Standard deviation of accuracy: {scores.std():.2f}")
```

OUTPUT:

Accuracy for each fold: [0.96666667 0.96666667 0.96666667 0.96666667 1.]

Mean accuracy: 0.97

Standard deviation of accuracy: 0.02

PROGRAM:

```
import numpy as np
from scipy import stats
import pandas as pd

alpha = 0.05

before_tuning = [70, 75, 80, 82, 78, 85, 83, 77, 74, 79]
after_tuning = [78, 80, 85, 88, 83, 90, 89, 82, 80, 86]

t_stat, p_value = stats.ttest_rel(before_tuning, after_tuning)

print(f"T-statistic: {t_stat}")
print(f"P-value: {p_value}")

if p_value < alpha:
    print("There is a significant difference between the before and after performance (reject H0).")
else:
    print("There is no significant difference between the before and after performance (fail to reject H0).")
```

OUTPUT:

T-statistic: -5.741978695179742

P-value: 0.00021542350254357943

There is a significant difference between the before and after performance (reject H_0).